

# Methods of Applied Stats , More INLA

Patrick Brown, University of Toronto and St Mike's

Sept to Dec 2020

## Example

```
> data("bacteria", package = "MASS")
> bacteria$infected = as.numeric(bacteria$y == "y")
> bacteria$weekFac = factor(bacteria$week)
> bacteria$trt = factor(bacteria$ap, levels = c("p", "a"),
+   labels = c("placebo", "active treatment"))
> bacteria[c(1, 7:10), c("ID", "weekFac", "trt", "infected")]
```

	ID	weekFac	trt	infected
1	X01	0	placebo	1
7	X02	6	active treatment	0
8	X02	11	active treatment	1
9	X03	0	active treatment	1
10	X03	2	active treatment	1

$$Y_{ij} \sim \text{Bern}(\rho_{ij})$$
  
$$\text{logit}(\rho_{ij}) = X_{ij}\beta + U_i$$
  
$$U_i \sim N(0, \sigma^2)$$

*person* (pointing to  $i$ )  
*time* (pointing to  $j$ )

Bacteria data with an individual-level random effect

```

> library(INLA)
> res = inla(infected ~ weekFac + trt +  $U_i \sim N(0, \sigma^2)$ ,  $\sigma$  has a (exponential) prior median of two.
+ f(ID, model='iid', prior = 'pc.prec', param = c(2, 0.5)),
+ control.fixed = list(
+   mean = 0, mean.intercept = 0,
+   prec = 100^(-2), prec.intercept = 100^(-2)),
+ family = 'binomial', data=bacteria,
+ control.inla = list(strategy='laplace', fast=FALSE))
> res$summary.fixed[, c(4, 3, 5)]

```

	0.5quant	0.025quant	0.975quant
(Intercept)	3.7763208	2.327576	5.725563272
weekFac2	0.1643283	-1.255522	1.636596884
weekFac4	-1.5439286	-2.916394	-0.317738083
weekFac6	-1.6769759	-3.075238	-0.435195431
weekFac11	-1.6882410	-3.069860	-0.470155061
trtactive treatment	-1.1593444	-2.549680	-0.007278442

$$\beta_0 \sim N(0, 100^2)$$

$$\begin{pmatrix} \beta_1 \\ \vdots \\ \beta_p \end{pmatrix} \sim N(\vec{0}, 100^2 I)$$

①  $\beta$ s are usually insensitive to the priors

② Unlike  $\sigma$ ,  $\beta$  can be negative.

③ Flat priors (e.g.,  $N(0, 100^2 I)$ ) makes sense.


④ Always use default prior for the fixed effect.

## Natural Scale

$$Y_{ij} \sim \text{Bern}[\text{odds}_{ij}/(1 + \text{odds}_{ij})]$$

$$\text{odds}_{ij} = \prod_p \exp(\beta_p)^{X_{ijp}} \exp(U_i)$$

$$U_i \sim N(0, \sigma^2)$$

-   $\exp(\beta_p) = \text{odds}_{ij}/\text{odds}_{k\ell}$
- if  $X_{ijp} = X_{k\ell p} + 1$
- and  $X_{ijq} = X_{k\ell q}, q \neq p$
- and  $U_i = U_k$

```
> resTable = rbind(  
+   'Baseline prob' = 1/(1+exp(-res$summary.fixed[1,c(4,3,5)])),  
+   exp(res$summary.fixed[-1,c(4,3,5)]),  
+   '$\\sigma$'=Pmisc::priorPost(res)$summary[,c(4,3,5)])  
> rownames(resTable) = gsub("Fac|trt", " ", rownames(resTable))  
> knitr::kable(resTable, digits=2, caption = 'Posterior medians and quantiles')
```

Table 1: Posterior medians and quantiles for the baseline probability, odds ratios, and standard deviation of the random effect.

viation of the random effect.

(Estimation)

	0.5quant	0.025quant	0.975quant	
Baseline prob	0.98	0.91	1.00	
odds ratio {	week 2	1.18	0.28	5.14
	week 4	0.21	0.05	0.73
	week 6	0.19	0.05	0.65
	week 11	0.18	0.05	0.62
	active treatment	0.31	0.08	0.99
$\sigma$	1.40	0.58	2.45	

- Looks like the active treatment is effective
- How do we interpret  $\sigma$ ?

## Natural scale for $\sigma$ ?

- $\exp(\sigma) = \text{odds}_{ij}/\text{odds}_{k\ell}$
- if  $U_i = U_k + \sigma$ ,  $X_{ijq} = X_{k\ell q}$
- i.e.  $U_k = 0$ 
  - $k$  is typical
- and  $U_i = \sigma$ 
  - $i$  is 1 SD more sickly than is typical

- $\sigma = 1.4$ ,  $\exp(\sigma) = 4.1$
- that's very big!
- Should we report  $\exp(\sigma)$  instead of  $\sigma$ ?
- nobody ever does, so it would be confusing.
- insisting on  $\sigma$  instead of  $\sigma^2$  or  $1/\sigma^2$  is as much as I'll challenge the standard practice for now.

$$Y_{ij} \sim \text{Bern}[\text{odds}_{ij}/(1 + \text{odds}_{ij})]$$

$$\text{odds}_{ij} = \prod_p \exp(\beta_p)^{X_{ijp}} \exp(U_i)$$

$$U_i \sim \text{N}(0, \sigma^2)$$

## Interpreting $\sigma$

- suppose  $U_k = 0$  and  $U_i = \sigma$
- $k$  is in the 50th percentile of sickness (Typical)
- $i$  is the  $100 * \Phi(1) = 100 * \text{pnorm}(1) = 84$ th percentile

### Inter-Quartile Range

- $k$  is the 25th percentile,  $i$  is the 75th.
- $U_i = \Phi^{-1}(0.75)\sigma = \text{qnorm}(0.75) \sigma \approx 0.67\sigma$
- $U_k = \Phi^{-1}(0.25)\sigma \approx -0.67\sigma$
- $\text{odds}_{ij} / \text{odds}_{kl} = \exp\{[\Phi^{-1}(0.75) - \Phi^{-1}(0.25)]\sigma\} \approx \exp(1.3\sigma)$
- $\text{IQR} = \exp(1.3 \cdot 1.4) = 6.6$
- huge!  $\text{IQR} \cdot \hat{\sigma}$

$$\frac{\exp(U_i)}{\exp(U_k)} = \frac{\exp(0.67\sigma)}{\exp(-0.67\sigma)}$$

Person in the 75th percentile of sickness has 6.6 times the odds of being sick as the person in the 25th.

## In reverse

- Suppose we don't know the results and we're setting a prior for  $\sigma$
- Is 0.5 a sensible prior median for  $\sigma$ ?
- The IQR for individual-level risk would be  $\exp(1.3 \cdot 0.5) = 1.9$
- that's fairly big.
- how about  $\sigma = 4$ ?  $\exp(1.3 \cdot 4) = 180$
- how about  $\sigma = 0.2$ ?  $\exp(1.3 \cdot 0.2) = 1.3$
- $\sigma = 0.2$  is about right.
- Suppose I want my prior median for  $\sigma$  to correspond to an IQR of 1.5
- ...  $\exp(1.3 \cdot \sigma) = 1.5$ , or  $\sigma = \log(1.5)/1.3 \approx 0.31$
- ... `prior = 'pc.prec', param = c(log(1.5)/1.3, 0.5)`
- or I want 1 sd to double the odds of infection
- ... so  $\sigma = \log(2) \approx 0.69$
- ... `prior = 'pc.prec', param = c(log(2), 0.5)`



## In summary

- we must set priors for  $\sigma$
- unless you have a reason to do otherwise, use an Exponential prior (or `pc.prec`)
- set the prior median, unless you have some reason to set the mean or rate or 95% quantile.
- posteriors for  $\sigma$  tend to be fairly insensitive to the hyperparameter of the Exponential.
- but you do need to justify the hyperparameter
- set  $pr(\sigma < \log(2)) = 0.5$ , prior median means 1 SD doubles odds *Sensible.*
- or  $pr(\sigma < \log(1.5)/1.3) = 0.5$ , the odds ratio of the IQR is a 50% increase
- or replace the 0.25 and 0.75 above with some other numbers (0.025 and 0.975?).

## Predicted values

Let's fit a treatment by time interaction model

```
> resGlm = glm(infected ~ weekFac * trt, family = binomial(link = "logit")  
+ data = bacteria)  
> summary(resGlm)$coef
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	2.25129180	0.7433919	3.02840495	0.0024584
weekFac2	0.69314717	1.2669470	0.54710037	0.5843097
weekFac4	-0.99852883	0.9349118	-1.06804599	0.2854997
weekFac6	0.52129692	1.2708682	0.41018961	0.6816668
weekFac11	-0.86499744	0.9301245	-0.92998029	0.3523812
trtactive treatment	-0.09180755	0.9614710	-0.09548655	0.9239283
weekFac2:trtactive treatment	-0.90672127	1.5355461	-0.59048781	0.5548636
weekFac4:trtactive treatment	-0.27365222	1.2031358	-0.22744916	0.8200744
weekFac6:trtactive treatment	-2.41841691	1.4709942	-1.64406967	0.1001618
weekFac11:trtactive treatment	-0.60133963	1.1934934	-0.50384832	0.6143679

## Another model

```
> resGlm2 = glm(infected ~ weekFac:trt, family = binomial(link = "logit"),  
+   data = bacteria)  
> summary(resGlm2)$coef
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	0.6931472	0.4330127	1.6007548	0.10943123
weekFac0:trtplacebo	1.5581446	0.8603090	1.8111453	0.07011836
weekFac2:trtplacebo	2.2512918	1.1135633	2.0217008	0.04320727
weekFac4:trtplacebo	0.5596158	0.7133923	0.7844433	0.43278006
weekFac6:trtplacebo	2.0794415	1.1180225	1.8599281	0.06289569
weekFac11:trtplacebo	0.6931472	0.7071068	0.9802581	0.32695871
weekFac0:trtactive treatment	1.4663371	0.7478602	1.9607101	0.04991284
weekFac2:trtactive treatment	1.2527630	0.7539578	1.6615823	0.09659655
weekFac4:trtactive treatment	0.1941560	0.6238435	0.3112255	0.75562918
weekFac6:trtactive treatment	-0.4307829	0.6036746	-0.7136011	0.47547385

the two models are the same, coefficients of one are linear combinations of the other

## Predicted

```
> newx = expand.grid(  
+   trt = levels(bacteria$trt),  
+   weekFac = levels(bacteria$weekFac))  
> newx
```

	trt	weekFac
1	placebo	0
2	active treatment	0
3	placebo	2
4	active treatment	2
5	placebo	4
6	active treatment	4
7	placebo	6
8	active treatment	6
9	placebo	11
10	active treatment	11

- suppose I want probability of being infected
- for each treatment and each week
- create a data frame newx to make predictions from

```
> myPred = as.data.frame(predict(  
+   resGlm, newx, se.fit=TRUE))  
> myPred[1:3,]
```

	fit	se.fit	residual.scale
1	2.251292	0.7433919	1
2	2.159484	0.6097498	1
3	2.944439	1.0259255	1

$X\beta \rightarrow \log \text{ odds.}$

## Probabilities

### convert to confidence intervals

```
> theCiMat = Pmisc::ciMat(0.95)
> theCiMat
```

	est	2.5	97.5
Estimate	1	1.000000	1.000000
Std. Error	0	-1.959964	1.959964

```
> myPredCI =
+   as.matrix(myPred[,1:2]) %*%
+   theCiMat
> myPredCI[1:3,] 95% CI for log odds
```


	est	2.5	97.5
1	2.251292	0.7942704	3.708313
2	2.159484	0.9643965	3.354572
3	2.944439	0.9336619	4.955216

### convert to probabilities

```
> myPredOdds = exp(myPredCI)
> myPredProb = myPredOdds /
+   (1+myPredOdds)
> myPredProb[1:3,]
```

	est	2.5	97.5
1	0.9047619	0.6887475	0.9760679
2	0.8965517	0.7240012	0.9662542
3	0.9500000	0.7178176	0.9930027

## A plot



*Estimated*

Figure: ~~Predicted~~ probabilities of infection over time by treatment group, with 95pct CI.

## The Code

```
> Scol = c('active treatment' = 'red', placebo = 'black')
> weekShift = as.numeric(newx$weekFac) +
+   (newx$trt=='active treatment')/5
> plot(weekShift, myPredProb[,1], pch = 15,
+   ylim = range(myPredProb), bty='n',
+   col = Scol[as.character(newx$trt)], xaxt='n',
+   xlab='week', ylab='prob')
      ↳ no axis
> segments(weekShift, myPredProb[,2], y1 = myPredProb[,3],
+   col = Scol[as.character(newx$trt)])
> axis(1, 1:nlevels(newx$weekFac), levels(newx$weekFac))
> legend('bottomleft', col=Scol, pch=15, pt.cex=2,
+   legend=names(Scol), bty='n')
```

## Predictions with INLA

This is, unfortunately, a bit of work

```
> forLincombs = do.call(INLA::inla.make.lincombs,  
+   as.data.frame(model.matrix( ~ weekFac*trt,  
+     data=newx)))  
  
> res2 = inla(infected ~ weekFac * trt +  
+   f(ID, model='iid', prior = 'pc.prec', param = c(log(2)2, 0.5))),  
+   control.fixed = list(  
+     mean = 0, mean.intercept = 0,  
+     prec = 100-2, prec.intercept = 100-2),  
+   lincomb = forLincombs,  
+   family = 'binomial', data=bacteria,  
+   control.compute = list(config=TRUE),  
+   control.inla = list(strategy='laplace', fast=FALSE))
```



## What we get

```
> cbind(newx, res2$summary.lincomb.derived[,c(5,4,6)])[1:7,]
```

	trt	week	Fac	0.5quant	0.025quant	0.975quant
lc01	placebo	0	3.232779	1.4125734	5.445617	
lc02	active treatment	0	2.898449	1.4359891	4.584478	
lc03	placebo	2	4.385277	2.0110292	7.094482	
lc04	active treatment	2	2.747086	1.2495438	4.474118	
lc05	placebo	4	2.017296	0.5780177	3.769063	
lc06	active treatment	4	1.558050	0.4001076	2.941508	
lc07	placebo	6	3.895301	1.5651737	6.475668	

- these are posterior quantiles of  $\tilde{X}\beta$
- where the 10 rows of  $\tilde{X}$  are the different treatment/week combinations
- If  $U_i = 0$  then  $\text{logit}(\rho_{ij}) = X_{ij}\beta$
- We've produced posterior quantiles for 'typical' people with  $U_i = 0$

## A plot




Figure: Posterior medians of probabilities of infection over time by treatment group, with 95pct CI.

## The code

```
> predOddsInla = exp(res2$summary.lincomb.derived[,c(5,4,6)])
> predProbInla = predOddsInla/(1+predOddsInla)
> plot(weekShift, predProbInla[,1], pch = 15,
+      ylim = range(predProbInla), bty='n',
+      col = Scol[as.character(newx$trt)], xaxt='n', xlab='week', ylab='prob')
> segments(weekShift, predProbInla[,2], y1 = predProbInla[,3],
+          col = Scol[as.character(newx$trt)])
> axis(1, 1:nlevels(newx$weekFac), levels(newx$weekFac))
> legend('bottomleft', col=Scol, pch=15, pt.cex=2,
+        legend=names(Scol), bty='n')
```

## In summary

- it doesn't matter whether you do `trt*week` or `trt:week` or `0 + trt*week`
- ... if you use `predict` or `lincomb` to get the values you want
- so, get into the habit of using them!
- and make nice graphs

## Posterior samples


- Joint posterior

$$\pi(\beta, U, \sigma|Y)$$

- Sample from this distribution?
- useful for calculating nonlinear things

```
> mySample =  
+   inla.posterior.sample(  
+     1000, res2)  
> length(mySample)  
[1] 1000  
> names(mySample[[1]])  
[1] "hyperpar" "latent"   "logd"
```

```
> sampleW = do.call(cbind,  
+   Biobase::subListExtract(  
+     mySample, 'latent'))  
> hist(sampleW['(Intercept):1',],  
+   prob=TRUE, main='', breaks=30, xlab='inte  
> lines(  
+   res2$marginals.fixed$('Intercept)', col=''
```



## How it works

- INLA approximates  $[\sigma|Y]$  discretely

- at values  $\sigma^{(1)} \dots \sigma^{(K)}$  — *configs.*

- first sample  $[\sigma^{(k)}|Y]$

- then sample  $[W|Y, \sigma^{(k)}]$  → *Normal approximation*

```
> mySample[[1]]$hyperpar
```

Precision for ID

0.2518642

```
> signif(sort(1/sqrt(exp(Biobase::subL
```

```
+ "theta", simplify = TRUE))))), 2)
```

[1] 0.29 0.36 0.44 0.55 0.68

[6] 0.85 1.00 1.30 1.60 2.00

[11] 2.50 3.10 3.80


```
> sampleTheta = Biobase::subListExtrac
```

```
+ mySample, 'hyperpar', simplify=T
```

```
> hist(1/sqrt(sampleTheta),
```

```
+ prob=TRUE, main='', breaks=50, xla
```

```
> lines(Pmisc::priorPost(res2)$sd$post
```



## some code

```
> sampleTheta = Biobase::subListExtract(  
+   mySample, 'hyperpar', simplify=TRUE)  
> hist(1/sqrt(sampleTheta),  
+   prob=TRUE, main='', breaks=50, xlab=expression(sigma))  
> lines(Pmisc::priorPost(res2)$sd$posterior, col='red', lwd=3)
```

## an alternative to linear combinations

```
> sampleW = do.call(cbind, Biobase::subListExtract(mySample,  
+   "latent"))  
> dim(sampleW)  
  
[1] 280 1000  
  
> # rownames(sampleW)  
> sampleBeta = t(sampleW[grepl("Intercept|weekFac|trt", rownames(sampleW)),  
+   ])  
> colnames(sampleBeta) = gsub(":1$", "", colnames(sampleBeta))  
> signif(sampleBeta[1:3, 1:4], 4)
```

	(Intercept)	weekFac2	weekFac4	weekFac6
sample:1	3.767	1.104	-1.402	-0.1067
sample:2	4.968	-1.444	-2.801	1.1110
sample:3	4.493	1.671	-2.986	-0.3771



We have samples of  $[\beta|Y]$  and want samples of  $[\tilde{X}\beta|Y]$ .

```
> newx2 = model.matrix(~weekFac * trt, newx)
```

```
> newx2[1:3, 1:4]
```

	(Intercept)	weekFac2	weekFac4	weekFac6
1	1	0	0	0
2	1	0	0	0
3	1	1	0	0

}  $\tilde{X}$ .

```
> sampleReparam = tcrossprod(sampleBeta[, colnames(newx2)],  
+ newx2)
```

```
> dim(sampleReparam)
```

```
[1] 1000  10
```

```
> colnames(sampleReparam) = colnames(newx2)
```

```
> sampleNatural = exp(sampleReparam)/(1 + exp(sampleReparam))
```

```
> t(signif(apply(sampleNatural, 2, quantile, prob = c(0.5,  
+ 0.025, 0.975))), 3))
```

	50%	2.5%	97.5%
(Intercept)	0.959	0.817	0.996
weekFac2	0.948	0.819	0.989
weekFac4	0.987	0.889	0.999
weekFac6	0.940	0.769	0.987
weekFac11	0.882	0.650	0.976
trtactive treatment	0.826	0.605	0.952
weekFac2:trtactive treatment	0.979	0.843	0.998
weekFac4:trtactive treatment	0.572	0.291	0.794
weekFac6:trtactive treatment	0.898	0.681	0.982
weekFac11:trtactive treatment	0.761	0.517	0.913

## References I